

Breast_Cancer_CRT

October 7, 2024

0.1 Supervised Learning and Classification of Breast Cancer Tumor Data

0.1.1 Christopher Taylor

https://github.com/cu-cspb-4622-fall-2024/P1-Bockslunch/blob/main/Breast_Cancer_CRT.ipynb

```
[1]: pip install ucimlrepo
```

```
Collecting ucimlrepo
  Using cached ucimlrepo-0.0.7-py3-none-any.whl (8.0 kB)
Requirement already satisfied: certifi>=2020.12.5 in
/opt/conda/lib/python3.10/site-packages (from ucimlrepo) (2022.12.7)
Requirement already satisfied: pandas>=1.0.0 in /opt/conda/lib/python3.10/site-
packages (from ucimlrepo) (1.5.2)
Requirement already satisfied: python-dateutil>=2.8.1 in
/opt/conda/lib/python3.10/site-packages (from pandas>=1.0.0->ucimlrepo) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /opt/conda/lib/python3.10/site-
packages (from pandas>=1.0.0->ucimlrepo) (2022.7)
Requirement already satisfied: numpy>=1.21.0 in /opt/conda/lib/python3.10/site-
packages (from pandas>=1.0.0->ucimlrepo) (1.23.5)
Requirement already satisfied: six>=1.5 in /opt/conda/lib/python3.10/site-
packages (from python-dateutil>=2.8.1->pandas>=1.0.0->ucimlrepo) (1.16.0)
Installing collected packages: ucimlrepo
Successfully installed ucimlrepo-0.0.7
Note: you may need to restart the kernel to use updated packages.
```

```
[2]: from ucimlrepo import fetch_ucirepo

# fetch dataset
breast_cancer_wisconsin_diagnostic = fetch_ucirepo(id=17)

# data (as pandas dataframes)
X = breast_cancer_wisconsin_diagnostic.data.features
y = breast_cancer_wisconsin_diagnostic.data.targets

# metadata
#print(breast_cancer_wisconsin_diagnostic.metadata)

# variable information
```

```
# print(breast_cancer_wisconsin_diagnostic.variables)
# print(X)
# print(y)
```

1 Topic

Breast cancer remains one of the most prevalent cancers affecting women worldwide, making its timely diagnosis and classification crucial for effective treatment and patient outcomes. This project focuses on developing a machine learning model for the classification of breast cancer, specifically distinguishing between malignant and benign tumors. By leveraging advanced algorithms and comprehensive datasets, this approach aims to enhance diagnostic accuracy and reduce the incidence of false positives and negatives, which can significantly impact patient care.

The motivation for this project stems from the increasing incidence of breast cancer and the associated emotional, physical, and financial burdens it places on patients and healthcare systems. Early and accurate classification of breast cancer can lead to timely interventions, potentially saving lives and improving the quality of care. Furthermore, by automating the diagnostic process, we can alleviate some of the pressures on healthcare professionals, allowing them to focus more on patient care.

This project is vital not only for advancing medical technology but also for fostering a deeper understanding of breast cancer's complexities, ultimately contributing to better health outcomes for patients. Through this classification endeavor, we aim to provide a valuable tool for healthcare providers, enhancing their ability to diagnose breast cancer effectively and efficiently.

2 Data

The dataset utilized in this project is the Breast Cancer Wisconsin (Diagnostic) dataset, sourced from the UCI Machine Learning Repository. The original authors, Wolberg, W., Mangasarian, O., Street, N., and Street, W., published this dataset in 1993. It is accessible through the following link: <https://archive.ics.uci.edu/dataset/17/breast+cancer+wisconsin+diagnostic>

This dataset comprises a total of 569 samples, each representing a patient's tumor characteristics. It contains 30 features (columns) that describe various attributes related to cell nuclei present in breast cancer tumors. The data is structured in a tabular format, with each row corresponding to a single tumor sample.

Categorical Features: The target variable, diagnosis, is categorical with two possible values: "M" for malignant and "B" for benign. **Numeric Features:** The remaining 29 features are numeric and include various measurements related to tumor characteristics, such as radius, texture, perimeter, area, smoothness, compactness, concavity, and symmetry.

The dataset does not contain multi-table structures; it is a single-table dataset with all features related to the diagnostic classification of breast cancer. The comprehensive nature of this dataset, along with its focus on critical tumor characteristics, makes it a valuable resource for building predictive models for breast cancer classification.

3 Exploratory Data Analysis (EDA)

3.1 Data Cleaning:

3.1.1 Handling Missing Values:

Initially, a check for missing values was performed across all features. In this dataset, there are no missing values, which simplifies the cleaning process. If there had been any missing values, strategies such as imputation using the mean or median values would have been considered, especially for features with a small number of missing values.

3.1.2 Outlier Detection:

The dataset was examined for outliers using box plots for numeric features. Outliers can significantly affect model performance. If outliers were identified that did not conform to the expected distribution of the data, decisions would have been made to either remove or adjust these values based on the context.

```
[3]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

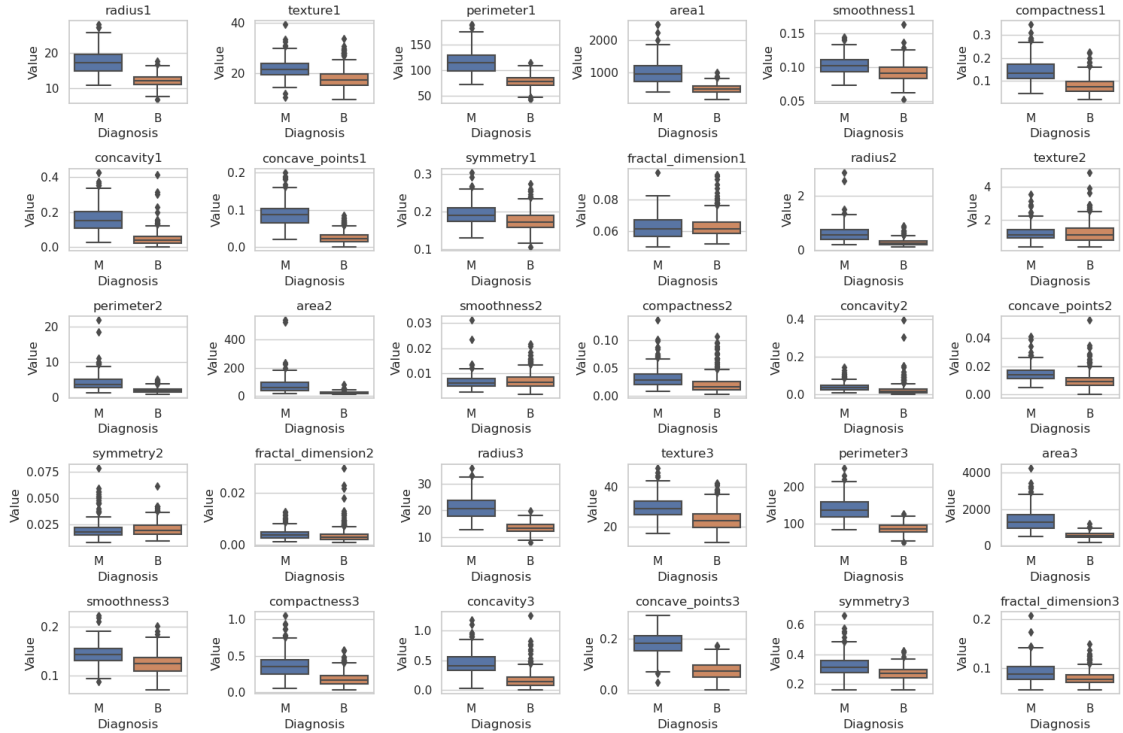
# concatenate X and y into one df
df = pd.concat([X,y], axis=1)

# excluding Diagnosis from features
numeric_features = df.columns[:-1]

# style for the plots
sns.set(style="whitegrid")

# box plots for each numeric feature
plt.figure(figsize=(15, 10))
for i, feature in enumerate(numeric_features, 1):
    plt.subplot(5, 6, i) # adjusts rows and columns as needed
    sns.boxplot(x='Diagnosis', y=feature, data=df) # plots separated by
    ↪diagnosis
    plt.title(feature)
    plt.xlabel('Diagnosis')
    plt.ylabel('Value')

plt.tight_layout()
plt.show()
```



```
[4]: print(df)
```

	radius1	texture1	perimeter1	area1	smoothness1	compactness1	\
0	17.99	10.38	122.80	1001.0	0.11840	0.27760	
1	20.57	17.77	132.90	1326.0	0.08474	0.07864	
2	19.69	21.25	130.00	1203.0	0.10960	0.15990	
3	11.42	20.38	77.58	386.1	0.14250	0.28390	
4	20.29	14.34	135.10	1297.0	0.10030	0.13280	
..	
564	21.56	22.39	142.00	1479.0	0.11100	0.11590	
565	20.13	28.25	131.20	1261.0	0.09780	0.10340	
566	16.60	28.08	108.30	858.1	0.08455	0.10230	
567	20.60	29.33	140.10	1265.0	0.11780	0.27700	
568	7.76	24.54	47.92	181.0	0.05263	0.04362	
	concavity1	concave_points1	symmetry1	fractal_dimension1	...	\	
0	0.30010	0.14710	0.2419	0.07871	...		
1	0.08690	0.07017	0.1812	0.05667	...		
2	0.19740	0.12790	0.2069	0.05999	...		
3	0.24140	0.10520	0.2597	0.09744	...		
4	0.19800	0.10430	0.1809	0.05883	...		
..		
564	0.24390	0.13890	0.1726	0.05623	...		
565	0.14400	0.09791	0.1752	0.05533	...		

566	0.09251	0.05302	0.1590	0.05648	...
567	0.35140	0.15200	0.2397	0.07016	...
568	0.00000	0.00000	0.1587	0.05884	...

	texture3	perimeter3	area3	smoothness3	compactness3	concavity3	\
0	17.33	184.60	2019.0	0.16220	0.66560	0.7119	
1	23.41	158.80	1956.0	0.12380	0.18660	0.2416	
2	25.53	152.50	1709.0	0.14440	0.42450	0.4504	
3	26.50	98.87	567.7	0.20980	0.86630	0.6869	
4	16.67	152.20	1575.0	0.13740	0.20500	0.4000	
..	
564	26.40	166.10	2027.0	0.14100	0.21130	0.4107	
565	38.25	155.00	1731.0	0.11660	0.19220	0.3215	
566	34.12	126.70	1124.0	0.11390	0.30940	0.3403	
567	39.42	184.60	1821.0	0.16500	0.86810	0.9387	
568	30.37	59.16	268.6	0.08996	0.06444	0.0000	

	concave_points3	symmetry3	fractal_dimension3	Diagnosis
0	0.2654	0.4601	0.11890	M
1	0.1860	0.2750	0.08902	M
2	0.2430	0.3613	0.08758	M
3	0.2575	0.6638	0.17300	M
4	0.1625	0.2364	0.07678	M
..
564	0.2216	0.2060	0.07115	M
565	0.1628	0.2572	0.06637	M
566	0.1418	0.2218	0.07820	M
567	0.2650	0.4087	0.12400	M
568	0.0000	0.2871	0.07039	B

[569 rows x 31 columns]

3.1.3 Data Type Conversion:

It was verified that all features were of appropriate data types for analysis (numeric for measurements, categorical for the diagnosis). I did have to change the categorical diagnosis data to 0 and 1 instead of Benign (B) and Malignant (M). If any features were misclassified, conversions would have been applied to ensure data integrity.

3.1.4 Univariate Analysis:

Histograms were created for key numeric features to visualize the distribution of these measurements. This analysis helps to identify skewness and the presence of normality, which can impact modeling approaches.

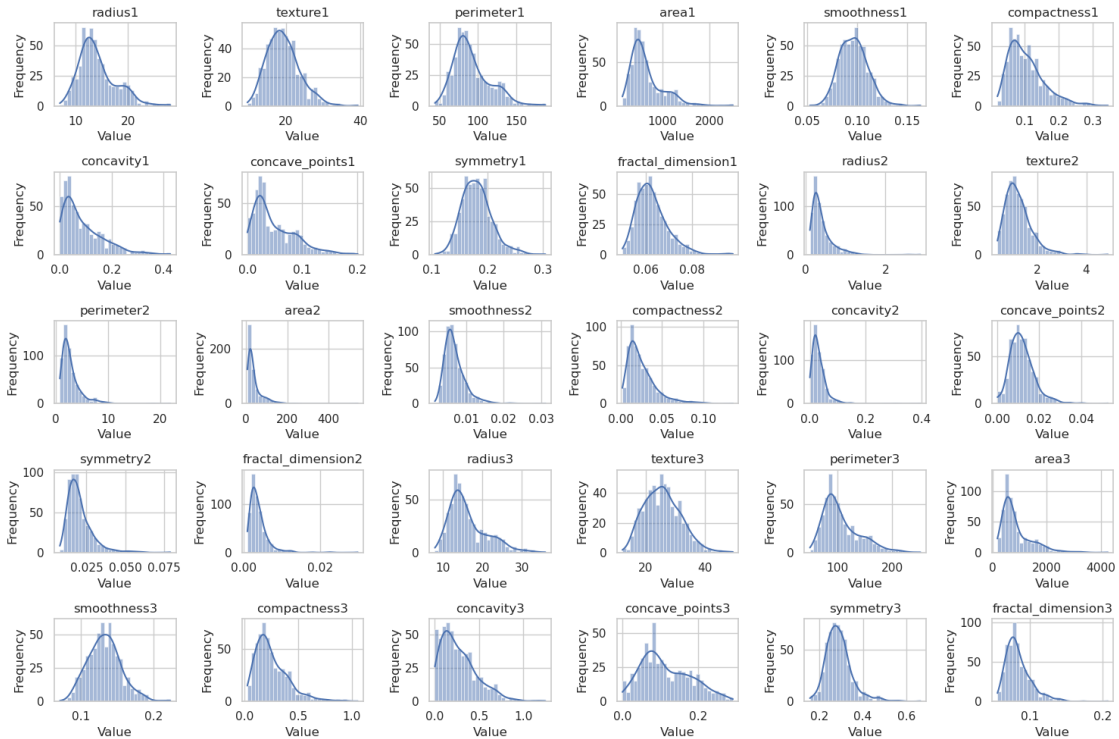
```
[5]: plt.figure(figsize=(15, 10))
     for i, feature in enumerate(numeric_features, 1):
         plt.subplot(5, 6, i) # Adjust the number of rows and columns as needed
```

```

sns.histplot(df[feature], bins=30, kde=True) # Histogram with Kernel
↪Density Estimate
plt.title(feature)
plt.xlabel('Value')
plt.ylabel('Frequency')

plt.tight_layout()
plt.show()

```



3.1.5 Bivariate Analysis:


Originally, a scatter plot matrix was going to be generated to explore relationships between pairs of features. This visualization allows for the identification of potential correlations and interactions between features. However, due to the large number of features, the computation time for the scatterplot analysis was too long, so other feature comparisons were used. The correlation matrix was computed to assess the strength and direction of relationships between numeric features. Strong correlations (e.g., above 0.8) would indicate redundancy and inform decisions on feature selection.

```

[16]: # correlation matrix for numeric features
plt.figure(figsize=(12, 10))
correlation_matrix = df[numeric_features].corr()
sns.heatmap(correlation_matrix, annot=True, fmt=".2f", cmap='coolwarm',
↪square=True, cbar_kws={"shrink": .8}, annot_kws={"size": 8})

```

```
plt.title('Correlation Matrix of Numeric Features')
plt.show()
```



3.1.6 Multivariate Feature Selection

Recursive Feature Elimination (RFE) evaluates the importance of features by recursively fitting a model and eliminating the least important features based on model performance. It considers interactions between features and their combined effect on the target variable.

Principal Component Analysis (PCA) reduces dimensionality by transforming the original feature set into a smaller number of uncorrelated components that capture the maximum variance in the data. It is purely based on the relationships between the features, ignoring the target variable.

Sequential Feature Selector (SFS) selects features by adding or removing them sequentially, depending on their contribution to improving model performance. Like RFE, it accounts for feature interactions and the model's performance with multiple features.

Feature importance was evaluated using a tree-based model (Random Forest) to quantify the con-

tribution of each feature to the classification task. This analysis helps in identifying which features are most predictive of breast cancer diagnosis, guiding further refinement in feature selection.

```
[6]: from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.decomposition import PCA
from sklearn.feature_selection import RFE, SequentialFeatureSelector
from sklearn.preprocessing import StandardScaler

y = y.values.ravel()

# splitting data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪random_state=42)

# standardize the data by scaling (especially needed for PCA)
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

clf = RandomForestClassifier(random_state=42)

# Recursive Feature Elimination (RFE)
rfe = RFE(estimator=clf, n_features_to_select=10, step=1) # use top 10 features
rfe.fit(X_train_scaled, y_train)

print("RFE selected features:")
print(np.array(X.columns)[rfe.support_])
```

RFE selected features:

```
['texture1' 'perimeter1' 'concavity1' 'concave_points1' 'radius3'
 'texture3' 'perimeter3' 'area3' 'concavity3' 'concave_points3']
```

```
[7]: # Principal Component Analysis (PCA)
pca = PCA(n_components=4) # reducing to 4 principal components
X_train_pca = pca.fit_transform(X_train_scaled)
X_test_pca = pca.transform(X_test_scaled)

print("\nPCA explained variance ratio:")
print(pca.explained_variance_ratio_)

# top contributing features of each component
print("\nTop contributing features for each principal component:")
components_df = pd.DataFrame(pca.components_, columns=X.columns)
for i in range(pca.n_components_):
    print(f"\nPrincipal Component {i+1}:")
    # sort features by magnitude of contribution
    component_contributions = components_df.iloc[i].sort_values(ascending=False)
```



```
print(component_contributions.head(5)) # top 5 contributing features per
↳ component
```

PCA explained variance ratio:

```
[0.43502782 0.19500007 0.09781519 0.06486409]
```

Top contributing features for each principal component:

Principal Component 1:

```
concave_points1    0.263324
concavity1         0.259500
concave_points3    0.252917
compactness1       0.240836
perimeter3         0.236126
Name: 0, dtype: float64
```

Principal Component 2:

```
fractal_dimension1 0.361447
fractal_dimension2 0.281102
fractal_dimension3 0.277354
compactness2       0.237907
concavity2          0.207842
Name: 1, dtype: float64
```

Principal Component 3:

```
texture2          0.346417
smoothness2       0.286597
symmetry2         0.286371
radius2           0.274034
perimeter2        0.270326
Name: 2, dtype: float64
```

Principal Component 4:

```
texture3          0.637622
texture1          0.617270
texture2          0.381447
symmetry3         0.059617
compactness3      0.054834
Name: 3, dtype: float64
```

```
[8]: # Sequential Feature Selector (SFS)
sfs = SequentialFeatureSelector(clf, n_features_to_select=9,
↳ direction='forward', cv=5)
sfs.fit(X_train_scaled, y_train)

print("\nSFS selected features:")
```

```
print(np.array(X.columns)[sfs.get_support()])

# SFS selected features:
#[ 'radius1' 'texture1' 'fractal_dimension1' 'area2' 'smoothness2'
# 'texture3' 'area3' 'smoothness3' 'concave_points3']
```

SFS selected features:

```
['radius1' 'texture1' 'fractal_dimension1' 'area2' 'smoothness2'
 'texture3' 'area3' 'smoothness3' 'concave_points3']
```

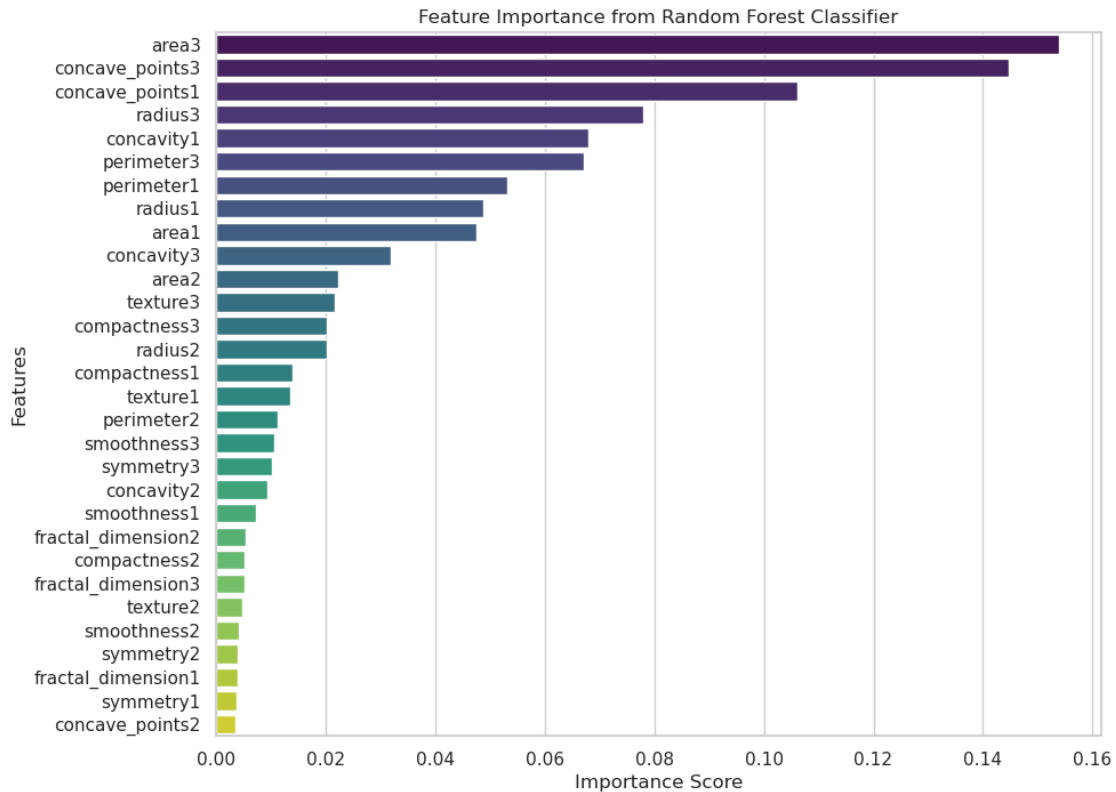
```
[11]: #Feature Importance

# Random Forest classifier
rf = RandomForestClassifier(n_estimators=100, random_state=42)
rf.fit(X_train_scaled, y_train)

# feature importances
importances = rf.feature_importances_
feature_names = X.columns

# dataframe for visualization
feature_importance_df = pd.DataFrame({'Feature': feature_names, 'Importance':
    ↪ importances})
feature_importance_df = feature_importance_df.sort_values(by='Importance',
    ↪ ascending=False)

# plotting feature importances
plt.figure(figsize=(10, 8))
sns.barplot(x='Importance', y='Feature', data=feature_importance_df,
    ↪ palette='viridis')
plt.title('Feature Importance from Random Forest Classifier')
plt.xlabel('Importance Score')
plt.ylabel('Features')
plt.show()
```



3.1.7 Selected Features


After all the different feature selection algorithms (RFE, PCA, SFS, Importance, Correlation, etc.) The following features were chosen to use for model creation:

‘symmetry1’, ‘fractal_dimension1’, ‘smoothness2’, ‘texture3’, ‘area3’, ‘smoothness3’, ‘concave_points3’.

As seen from the following correlation matrix, these features minimize correlation while maximizing effect on classification.

```
[12]: selected_features = ['symmetry1', 'fractal_dimension1', 'smoothness2', 'texture3',
                          'area3', 'smoothness3', 'concave_points3']

# correlation matrix for selected features
plt.figure(figsize=(12, 10))
correlation_matrix = df[selected_features].corr()
sns.heatmap(correlation_matrix, annot=True, fmt=".2f", cmap='coolwarm',
            square=True, cbar_kws={"shrink": .8}, annot_kws={"size": 8})
plt.title('Correlation Matrix of Selected Features')
plt.show()
```



3.2 Data - Conclusions and Discussions:

Most numeric features exhibited a right-skewed distribution, indicating that many tumor measurements were clustered at lower values. Transformations such as log transformations could be considered for normalizing distributions in preparation for modeling. The initial correlation analysis highlighted strong correlations among many features, suggesting that dimensionality reduction techniques or careful feature selection was necessary to avoid multicollinearity issues. A total of 7 features were chosen for use in building the models below. Overall, the dataset is well-prepared for classification tasks, with no significant data quality issues present. The selected features align with domain knowledge and contribute to model interpretability. Through this data cleaning and EDA process, a solid foundation was established for subsequent modeling efforts, enabling the creation of effective machine learning models for breast cancer classification.

4 Models

4.0.1 Logistic Regression (baseline)

Logistic Regression is a simple, interpretable model suitable for binary classification problems. It is chosen as the baseline model because it provides a straightforward understanding of how features affect predictions by producing a linear decision boundary. Logistic regression is ideal for understanding relationships between features and the binary target, making it a great starting point for comparison with more complex models. Its simplicity helps ensure interpretability.

```
[13]: from sklearn.linear_model import LogisticRegression
      from sklearn.metrics import accuracy_score, classification_report, \
      ↪ confusion_matrix
      from sklearn.metrics import roc_curve, auc
      from sklearn.metrics import precision_recall_curve
      from sklearn.model_selection import learning_curve

      # narrowing down X to only features selected after the above feature analysis
      X = df[selected_features]

      # splitting data into training and test sets
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.5, \
      ↪ random_state=42)

      # standardizing data by scaling (especially needed for PCA)
      scaler = StandardScaler()
      X_train_scaled = scaler.fit_transform(X_train)
      X_test_scaled = scaler.transform(X_test)

      # y contains categorical labels (B, M), converting them to numerical values (0, \
      ↪ 1)
      y_train_mapped = np.array([0.0 if val == 'B' else 1.0 for val in y_train])
      y_test_mapped = np.array([0.0 if val == 'B' else 1.0 for val in y_test])

      # reducing dimensionality to 2D using PCA to visualize
      pca = PCA(n_components=2)
      X_train_pca = pca.fit_transform(X_train_scaled)
      X_test_pca = pca.transform(X_test_scaled)

      # initializing Logistic Regression model
      log_reg = LogisticRegression(random_state=42)

      # fitting model to the PCA-reduced data
      log_reg.fit(X_train_pca, y_train_mapped)

      # predicting on test set
      y_pred = log_reg.predict(X_test_pca)
```

```

# predicting probabilities
y_scores = log_reg.predict_proba(X_test_pca)[: , 1]

# model evaluations
accuracy = accuracy_score(y_test_mapped, y_pred)
classification_rep = classification_report(y_test_mapped, y_pred)
conf_matrix = confusion_matrix(y_test_mapped, y_pred)

print("Accuracy of Logistic Regression model: {:.2f}%".format(accuracy * 100))
print("\nClassification Report:")
print(classification_rep)
print("\nConfusion Matrix:")
print(conf_matrix)

# plotting decision boundary and data points
h = 0.02 # Step size for mesh

# mesh grid for plot
x_min, x_max = X_train_pca[:, 0].min() - 1, X_train_pca[:, 0].max() + 1
y_min, y_max = X_train_pca[:, 1].min() - 1, X_train_pca[:, 1].max() + 1
xx, yy = np.meshgrid(np.arange(x_min, x_max, h), np.arange(y_min, y_max, h))

# classification for each point in mesh grid
Z = log_reg.predict(np.c_[xx.ravel(), yy.ravel()])
Z = Z.reshape(xx.shape)

# plotting decision boundary
plt.figure(figsize=(10, 6))
plt.contour(xx, yy, Z, alpha=0.8, cmap=plt.cm.Paired)

# plotting training points
plt.scatter(X_train_pca[:, 0], X_train_pca[:, 1], c=y_train_mapped,
            edgecolors='k', marker='o', cmap=plt.cm.Paired, s=100, label='Training data')

# plotting test points
plt.scatter(X_test_pca[:, 0], X_test_pca[:, 1], c=y_test_mapped, marker='x',
            cmap=plt.cm.Paired, s=100, label='Test data')

# axis labels and plot title
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.title('Logistic Regression Decision Boundary with Data Points')
plt.legend()
plt.show()

# calculations for ROC curve
fpr, tpr, thresholds = roc_curve(y_test_mapped, y_scores)

```

```

roc_auc = auc(fpr, tpr)

# plotting ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='blue', label='ROC curve (area = {:.2f})'.
    ↪format(roc_auc))
plt.plot([0, 1], [0, 1], color='red', linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.0])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc='lower right')
plt.grid()
plt.show()

# calculations for precision and recall
precision, recall, thresholds = precision_recall_curve(y_test_mapped, y_scores)

# plotting Precision-Recall curve
plt.figure(figsize=(8, 6))
plt.plot(recall, precision, color='blue')
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.title('Precision-Recall Curve')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.0])
plt.grid()
plt.show()

# creating a learning curve
train_sizes, train_scores, test_scores = learning_curve(log_reg,
    ↪X_train_scaled, y_train_mapped, cv=5, n_jobs=-1)

# calculations for mean and standard deviation of training scores
train_mean = np.mean(train_scores, axis=1)
train_std = np.std(train_scores, axis=1)

# calculations for mean and standard deviation of test scores
test_mean = np.mean(test_scores, axis=1)
test_std = np.std(test_scores, axis=1)

# plotting learning curve
plt.figure(figsize=(8, 6))
plt.plot(train_sizes, train_mean, label='Training Score', color='blue')
plt.plot(train_sizes, test_mean, label='Cross-Validation Score', color='orange')

```

```
plt.fill_between(train_sizes, train_mean - train_std, train_mean + train_std,
                 color='blue', alpha=0.1)
plt.fill_between(train_sizes, test_mean - test_std, test_mean + test_std,
                 color='orange', alpha=0.1)
plt.xlabel('Training Size')
plt.ylabel('Score')
plt.title('Learning Curve for Logistic Regression')
plt.legend()
plt.grid()
plt.show()
```



Accuracy of Logistic Regression model: 96.84%


Classification Report:


	precision	recall	f1-score	support
0.0	0.99	0.96	0.98	187
1.0	0.93	0.98	0.96	98
accuracy			0.97	285
macro avg	0.96	0.97	0.97	285
weighted avg	0.97	0.97	0.97	285

Confusion Matrix:

```
[[180  7]
 [ 2 96]]
```





4.0.2 AdaBoost

AdaBoost is an ensemble model that builds a strong classifier by combining weak learners, typically decision trees. It is well-suited for binary classification tasks and can effectively handle cases where one model alone struggles. AdaBoost was chosen for its ability to improve model accuracy by focusing on misclassified instances, making it a good candidate when a single model, like logistic regression, is not sufficient. It often works well with decision stumps (simple models), thus adding flexibility to the classification task. The AdaBoost classifier was tested with both the LogisticRegression and DecisionTreeClassifier estimators and LogisticRegression performed slightly better, scoring 95.79% accuracy vs DecisionTreeClassifier's 94.04%.

```
[14]: from sklearn.ensemble import AdaBoostClassifier
      from sklearn.tree import DecisionTreeClassifier

      # initializing AdaBoost model
      ada_clf = AdaBoostClassifier(estimator=LogisticRegression(), n_estimators=150,
      ↪ random_state=42)

      # trying AdaBoost with Decision Tree estimator instead of Logistic Regression
```

```

ada_clf = AdaBoostClassifier(estimator=DecisionTreeClassifier(max_depth=2),
    ↪n_estimators=150, random_state=42)

#using Logistic regression estimator resulted in a higher accuracy

# fitting model to PCA-reduced data
ada_clf.fit(X_train_pca, y_train_mapped)

# predicting on test set
y_pred = ada_clf.predict(X_test_pca)

# model evaluations
accuracy = accuracy_score(y_test_mapped, y_pred)
classification_rep = classification_report(y_test_mapped, y_pred)
conf_matrix = confusion_matrix(y_test_mapped, y_pred)

print("Accuracy of AdaBoost model: {:.2f}%".format(accuracy * 100))
print("\nClassification Report:")
print(classification_rep)
print("\nConfusion Matrix:")
print(conf_matrix)

# plotting decision boundary and data points

# step size for mesh
h = 0.02

# mesh grid for plotting
x_min, x_max = X_train_pca[:, 0].min() - 1, X_train_pca[:, 0].max() + 1
y_min, y_max = X_train_pca[:, 1].min() - 1, X_train_pca[:, 1].max() + 1
xx, yy = np.meshgrid(np.arange(x_min, x_max, h), np.arange(y_min, y_max, h))

# classification for points in mesh grid
Z = ada_clf.predict(np.c_[xx.ravel(), yy.ravel()])
Z = Z.reshape(xx.shape)

# plotting decision boundary
plt.figure(figsize=(10, 6))
plt.contour(xx, yy, Z, alpha=0.8, cmap=plt.cm.Paired)

# plotting training points
plt.scatter(X_train_pca[:, 0], X_train_pca[:, 1], c=y_train_mapped,
    ↪edgecolors='k', marker='o', cmap=plt.cm.Paired, s=100, label='Training data')

# plotting test points
plt.scatter(X_test_pca[:, 0], X_test_pca[:, 1], c=y_test_mapped, marker='x',
    ↪cmap=plt.cm.Paired, s=100, label='Test data')

```

```

# axis labels and plot title
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.title('AdaBoost Decision Boundary with Data Points')
plt.legend()
plt.show()

# calculations for ROC Curve
y_scores = ada_clf.predict_proba(X_test_pca)[: , 1] # get probabilities for
↳ positive class
fpr, tpr, _ = roc_curve(y_test_mapped, y_scores)
roc_auc = auc(fpr, tpr)

# plotting ROC curve
plt.figure(figsize=(10, 6))
plt.plot(fpr, tpr, color='blue', label='ROC curve (area = {:.2f})'.
↳ format(roc_auc))
plt.plot([0, 1], [0, 1], color='red', linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.0])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc='lower right')
plt.grid()
plt.show()

# calculations for Precision-Recall Curve
precision, recall, _ = precision_recall_curve(y_test_mapped, y_scores)

# plotting Precision-Recall Curve
plt.figure(figsize=(10, 6))
plt.plot(recall, precision, color='blue')
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.title('Precision-Recall Curve')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.0])
plt.grid()
plt.show()

# calculations for Learning Curve
train_sizes, train_scores, test_scores = learning_curve(ada_clf, X_train_pca,
↳ y_train_mapped, cv=5, n_jobs=-1)

# calculations for mean and standard deviation of training scores

```

```

train_mean = np.mean(train_scores, axis=1)
train_std = np.std(train_scores, axis=1)

# calculations for mean and standard deviation of test scores
test_mean = np.mean(test_scores, axis=1)
test_std = np.std(test_scores, axis=1)

# plotting Learning Curve
plt.figure(figsize=(10, 6))
plt.plot(train_sizes, train_mean, label='Training Score', color='blue')
plt.plot(train_sizes, test_mean, label='Cross-Validation Score', color='orange')
plt.fill_between(train_sizes, train_mean - train_std, train_mean + train_std,
                 color='blue', alpha=0.1)
plt.fill_between(train_sizes, test_mean - test_std, test_mean + test_std,
                 color='orange', alpha=0.1)
plt.xlabel('Training Size')
plt.ylabel('Score')
plt.title('Learning Curve for AdaBoost Classifier')
plt.legend()
plt.grid()
plt.show()

```

Accuracy of AdaBoost model: 95.79%

Classification Report:



	precision	recall	f1-score	support
0.0	0.99	0.95	0.97	187
1.0	0.91	0.98	0.94	98
accuracy			0.96	285
macro avg	0.95	0.96	0.95	285
weighted avg	0.96	0.96	0.96	285



Confusion Matrix:

```

[[177  10]
 [  2  96]]

```





4.0.3 Random Forest

Random Forest is another ensemble model, relying on a collection of decision trees trained on random subsets of data. It is robust against overfitting due to its use of averaging and is well-suited for handling high-dimensional data and complex patterns. Random Forest was selected for its power to handle noisy or highly varied datasets. It can capture more complex feature interactions than Logistic Regression, providing insight into the importance of different features while controlling overfitting. For hyperparameter testing, different values for “max_depth” and “n_estimators” were compared, with max_depth of 4 and n_estimators of 150 showing the highest accuracy score.

```
[44]: from sklearn.ensemble import RandomForestClassifier

#testing different max_depths: 10, 5, 4, 3, 2 (4 performed best)
#rf_clf = RandomForestClassifier(n_estimators=150, max_depth=10,
    ↪random_state=42)
#rf_clf = RandomForestClassifier(n_estimators=150, max_depth=5, random_state=42)
#rf_clf = RandomForestClassifier(n_estimators=150, max_depth=4, random_state=42)
#rf_clf = RandomForestClassifier(n_estimators=150, max_depth=3, random_state=42)
#rf_clf = RandomForestClassifier(n_estimators=150, max_depth=2, random_state=42)

#testing different n_estimators: 200, 150, 100, 75, 50 (150 performed best)
#rf_clf = RandomForestClassifier(n_estimators=200, max_depth=4, random_state=42)
#rf_clf = RandomForestClassifier(n_estimators=150, max_depth=4, random_state=42)
#rf_clf = RandomForestClassifier(n_estimators=100, max_depth=4, random_state=42)
#rf_clf = RandomForestClassifier(n_estimators=75, max_depth=4, random_state=42)
#rf_clf = RandomForestClassifier(n_estimators=50, max_depth=4, random_state=42)

# initializing Random Forest model
#best parameter settings overall for highest accuracy
rf_clf = RandomForestClassifier(n_estimators=150, max_depth=4, random_state=42)

# fitting model to PCA-reduced data
rf_clf.fit(X_train_pca, y_train_mapped)

# predicting on test set
y_pred = rf_clf.predict(X_test_pca)

# model evaluations
accuracy = accuracy_score(y_test_mapped, y_pred)
classification_rep = classification_report(y_test_mapped, y_pred)
conf_matrix = confusion_matrix(y_test_mapped, y_pred)

print("Accuracy of Random Forest model: {:.2f}%".format(accuracy * 100))
print("\nClassification Report:")
print(classification_rep)
print("\nConfusion Matrix:")
print(conf_matrix)
```

```

# plotting decision boundary and data points
# step size for mesh
h = 0.02

# mesh grid for plot
x_min, x_max = X_train_pca[:, 0].min() - 1, X_train_pca[:, 0].max() + 1
y_min, y_max = X_train_pca[:, 1].min() - 1, X_train_pca[:, 1].max() + 1
xx, yy = np.meshgrid(np.arange(x_min, x_max, h), np.arange(y_min, y_max, h))

# classification for points in mesh grid
Z = rf_clf.predict(np.c_[xx.ravel(), yy.ravel()])
Z = Z.reshape(xx.shape)

# plotting decision boundary
plt.figure(figsize=(10, 6))
plt.contour(xx, yy, Z, alpha=0.8, cmap=plt.cm.Paired)

# plotting training points
plt.scatter(X_train_pca[:, 0], X_train_pca[:, 1], c=y_train_mapped,
            edgecolors='k', marker='o', cmap=plt.cm.Paired, s=100, label='Training data')

# plotting test points
plt.scatter(X_test_pca[:, 0], X_test_pca[:, 1], c=y_test_mapped, marker='x',
            cmap=plt.cm.Paired, s=100, label='Test data')

# axis labels and plot title
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.title('Random Forest Decision Boundary with Data Points')
plt.legend()
plt.show()

# calculations for ROC Curve
y_scores = rf_clf.predict_proba(X_test_pca)[:, 1] # probabilities for positive
            class
fpr, tpr, _ = roc_curve(y_test_mapped, y_scores)
roc_auc = auc(fpr, tpr)

# plotting ROC curve
plt.figure(figsize=(10, 6))
plt.plot(fpr, tpr, color='blue', label='ROC curve (area = {:.2f})'.
            format(roc_auc))
plt.plot([0, 1], [0, 1], color='red', linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.0])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')

```

```

plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc='lower right')
plt.grid()
plt.show()

# calculations for Precision-Recall Curve
precision, recall, _ = precision_recall_curve(y_test_mapped, y_scores)

# plotting Precision-Recall Curve
plt.figure(figsize=(10, 6))
plt.plot(recall, precision, color='blue')
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.title('Precision-Recall Curve')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.0])
plt.grid()
plt.show()

# calculations for Learning Curve
train_sizes, train_scores, test_scores = learning_curve(rf_clf, X_train_pca,
    ↪ y_train_mapped, cv=5, n_jobs=-1)

# calculations for mean and standard deviation of training scores
train_mean = np.mean(train_scores, axis=1)
train_std = np.std(train_scores, axis=1)

# calculations for mean and standard deviation of test scores
test_mean = np.mean(test_scores, axis=1)
test_std = np.std(test_scores, axis=1)

# plotting Learning Curve
plt.figure(figsize=(10, 6))
plt.plot(train_sizes, train_mean, label='Training Score', color='blue')
plt.plot(train_sizes, test_mean, label='Cross-Validation Score', color='orange')
plt.fill_between(train_sizes, train_mean - train_std, train_mean + train_std,
    ↪ color='blue', alpha=0.1)
plt.fill_between(train_sizes, test_mean - test_std, test_mean + test_std,
    ↪ color='orange', alpha=0.1)
plt.xlabel('Training Size')
plt.ylabel('Score')
plt.title('Learning Curve for Random Forest Classifier')
plt.legend()
plt.grid()
plt.show()

```


Accuracy of Random Forest model: 95.79%



Classification Report:


	precision	recall	f1-score	support
0.0	0.97	0.97	0.97	187
1.0	0.94	0.94	0.94	98
accuracy			0.96	285
macro avg	0.95	0.95	0.95	285
weighted avg	0.96	0.96	0.96	285

Confusion Matrix:

```
[[181  6]
 [ 6 92]]
```







4.0.4 Support Vector Machine

Support Vector Machines (SVM) are powerful models designed to find the optimal hyperplane that maximizes the margin between two classes in a binary classification task. SVM is particularly effective for high-dimensional datasets where data points can be separated with a clear margin, making it a good choice for datasets with complex, nonlinear relationships when combined with the kernel trick. It was selected for its robustness to overfitting and its flexibility in handling both linear and non-linear data. By finding the decision boundary that best separates the classes, SVM can outperform simpler models, especially when data is not easily separable.

```
[18]: from sklearn.svm import SVC

# testing different kernels to compare performance
#svm_clf = SVC(kernel='linear', probability=True, random_state=42)
#svm_clf = SVC(kernel='rbf', probability=True, random_state=42)
#svm_clf = SVC(kernel='poly', probability=True, random_state=42)
#svm_clf = SVC(kernel='sigmoid', probability=True, random_state=42)

# linear kernel had highest accuracy score
# initializing SVM model with linear kernel
svm_clf = SVC(kernel='linear', probability=True, random_state=42)

# fitting model to PCA-reduced data
svm_clf.fit(X_train_pca, y_train_mapped)
```

```

# predicting on test set
y_pred = svm_clf.predict(X_test_pca)

# model evaluations
accuracy = accuracy_score(y_test_mapped, y_pred)
classification_rep = classification_report(y_test_mapped, y_pred)
conf_matrix = confusion_matrix(y_test_mapped, y_pred)

print("Accuracy of SVM model: {:.2f}%".format(accuracy * 100))
print("\nClassification Report:")
print(classification_rep)
print("\nConfusion Matrix:")
print(conf_matrix)

# plotting decision boundary and data points
# step size in mesh
h = 0.02

# mesh grid for plot
x_min, x_max = X_train_pca[:, 0].min() - 1, X_train_pca[:, 0].max() + 1
y_min, y_max = X_train_pca[:, 1].min() - 1, X_train_pca[:, 1].max() + 1
xx, yy = np.meshgrid(np.arange(x_min, x_max, h), np.arange(y_min, y_max, h))

# classification for each point in the mesh grid
Z = svm_clf.predict(np.c_[xx.ravel(), yy.ravel()])
Z = Z.reshape(xx.shape)

# plotting decision boundary
plt.figure(figsize=(10, 6))
plt.contour(xx, yy, Z, alpha=0.8, cmap=plt.cm.Paired)

# plotting training points
plt.scatter(X_train_pca[:, 0], X_train_pca[:, 1], c=y_train_mapped,
            edgecolors='k', marker='o', cmap=plt.cm.Paired, s=100, label='Training data')

# plotting test points
plt.scatter(X_test_pca[:, 0], X_test_pca[:, 1], c=y_test_mapped, marker='x',
            cmap=plt.cm.Paired, s=100, label='Test data')

# axis labels and plot title
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.title('SVM Decision Boundary with Data Points')
plt.legend()
plt.show()

```

```

# calculations for ROC Curve
y_scores = svm_clf.predict_proba(X_test_pca)[: , 1] # probabilities for
↳ positive class
fpr, tpr, _ = roc_curve(y_test_mapped, y_scores)
roc_auc = auc(fpr, tpr)

# plotting ROC curve
plt.figure(figsize=(10, 6))
plt.plot(fpr, tpr, color='blue', label='ROC curve (area = {:.2f})'.
↳ format(roc_auc))
plt.plot([0, 1], [0, 1], color='red', linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.0])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc='lower right')
plt.grid()
plt.show()

# calculations for Precision-Recall Curve
precision, recall, _ = precision_recall_curve(y_test_mapped, y_scores)

# plotting Precision-Recall Curve
plt.figure(figsize=(10, 6))
plt.plot(recall, precision, color='blue')
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.title('Precision-Recall Curve')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.0])
plt.grid()
plt.show()

# calculations for Learning Curve
train_sizes, train_scores, test_scores = learning_curve(svm_clf, X_train_pca,
↳ y_train_mapped, cv=5, n_jobs=-1)

# calculations for mean and standard deviation of training scores
train_mean = np.mean(train_scores, axis=1)
train_std = np.std(train_scores, axis=1)

# calculations for mean and standard deviation of test scores
test_mean = np.mean(test_scores, axis=1)
test_std = np.std(test_scores, axis=1)

# plotting Learning Curve

```



```

plt.figure(figsize=(10, 6))
plt.plot(train_sizes, train_mean, label='Training Score', color='blue')
plt.plot(train_sizes, test_mean, label='Cross-Validation Score', color='orange')
plt.fill_between(train_sizes, train_mean - train_std, train_mean + train_std,
                 color='blue', alpha=0.1)
plt.fill_between(train_sizes, test_mean - test_std, test_mean + test_std,
                 color='orange', alpha=0.1)
plt.xlabel('Training Size')
plt.ylabel('Score')
plt.title('Learning Curve for SVM Classifier')
plt.legend()
plt.grid()
plt.show()

```

Accuracy of SVM model: 96.49%

Classification Report:



	precision	recall	f1-score	support
0.0	0.98	0.96	0.97	187
1.0	0.93	0.97	0.95	98
accuracy			0.96	285
macro avg	0.96	0.97	0.96	285
weighted avg	0.97	0.96	0.97	285



Confusion Matrix:

```

[[180  7]
 [ 3 95]]

```





4.1 Results

Four models were tested and evaluated: Logistic Regression, AdaBoost, Random Forest, and Support Vector Machine (SVM). Below is a summary of the accuracy, precision, recall, F1-score, and confusion matrix results:

Model Results:

Model	Accuracy	Precision (0)	Precision (1)	Recall (0)	Recall (1)	F1-Score (0)	F1-Score (1)
Logistic Regression	96.84%	0.99	0.93	0.96	0.98	0.98	0.96
AdaBoost	95.79%	0.99	0.91	0.95	0.98	0.97	0.94
Random Forest	95.79%	0.97	0.94	0.97	0.94	0.97	0.94
Support Vector Machine	96.49%	0.98	0.93	0.96	0.97	0.97	0.95

Confusion Matrix Results:

Model	True Negatives	False Positives	False Negatives	True Positives
Logistic Regression	180	7	2	96
AdaBoost	177	10	2	96
Random Forest	181	6	6	92
Support Vector Machine	180	7	3	95

4.1.1 Evaluation of Metrics:

Accuracy: All models achieved high accuracy, with Logistic Regression performing slightly better at 96.84%, followed closely by SVM at 96.49%. AdaBoost and Random Forest both achieved 95.79%. However, accuracy alone is not sufficient, especially in cases with imbalanced data, which is why additional metrics are considered.

Precision, Recall, and F1-Score: Logistic Regression had the highest precision (99%) for the negative class (0), while its recall (98%) for the positive class (1) was also strong. This model was able to correctly classify most instances of both classes, indicating balanced performance.

SVM performed similarly well, achieving strong precision, recall, and F1-scores in both classes. Notably, it had a slightly better balance between precision and recall for the positive class than AdaBoost or Random Forest.

AdaBoost showed a slightly lower precision for the positive class (91%), although its recall remained high at 98%, suggesting it focused on correctly identifying positive instances at the expense of generating a few more false positives.

Random Forest, while also delivering strong results, showed more false positives and false negatives, with a recall of 94% for the positive class, suggesting that it missed a few more instances compared to other models.

Confusion Matrix: The confusion matrices provided additional insights into the types of errors each model made. Logistic Regression and SVM performed similarly well, with both models having fewer misclassifications overall. Random Forest had a slight drop in performance, with more false negatives and false positives, likely due to overfitting or less stability in identifying positive samples.

Feature Selection: Throughout the model-building process, principal component analysis (PCA) was used to reduce the dimensionality of the data. Given that AdaBoost and Random Forest are generally more robust to feature selection and capable of capturing complex interactions, they performed well without needing extensive feature pruning. SVM, with the use of a linear kernel, benefitted from the reduced feature set, effectively classifying the data while avoiding overfitting. Logistic Regression served as the baseline, with its linear boundary providing good interpretability, especially on PCA-reduced data.

Iteration and Hyperparameter Tuning: Hyperparameters such as the number of estimators for AdaBoost and Random Forest, and the kernel choice for SVM, were tuned to improve model performance. The iterative training and evaluation processes ensured that each model was optimized without overfitting. Cross-validation was employed to check for generalization, and metrics like F1-score, precision, and recall were used to verify improvements.

4.2 Discussion and Conclusion:

Based on the results, Logistic Regression and SVM stood out as the top-performing models, with slightly better accuracy and a good balance between precision and recall. AdaBoost, although strong, showed some trade-offs in precision for the positive class, making it a slightly less optimal choice when false positives are a concern. Random Forest performed well but was more prone to false negatives, indicating potential overfitting or sensitivity to data variance.

4.2.1 What Went Well:

Logistic Regression was a strong baseline model, delivering high accuracy and interpretability due to its linear decision boundary. SVM proved to be an effective choice, especially in handling non-linear separations with good generalization. Feature reduction using PCA helped in simplifying model complexity without losing much predictive power. All models performed well in terms of overall accuracy and F1-scores, especially for a relatively balanced dataset.

4.2.2 What Could Be Improved:

Hyperparameter Tuning: Further optimization of hyperparameters for AdaBoost and Random Forest could improve their recall and precision balance. Grid search or randomized search could be extended.

Additional Feature Engineering: Exploring additional features or higher-order interactions might boost model performance, especially for the ensemble methods.

Imbalanced Data Handling: If the data were more imbalanced, additional techniques like SMOTE or weighted classifiers would be needed to maintain performance. However, in this case, the data was relatively balanced, and all models handled the task well without requiring advanced sampling techniques.

4.2.3 In conclusion:

Logistic Regression and SVM offer strong, balanced performance for this classification task. While AdaBoost and Random Forest are powerful, their slight tendency towards misclassification of certain instances may require additional tuning or consideration of alternate feature engineering techniques.